

```
#Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import datetime
```

```
#importing the dataset
df= pd.read_csv("/content/uber.csv")
```

1. Pre-process the dataset.

```
df.head()
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.723217
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.750325
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.772647
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.803349
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.761247



To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```
RangeIndex: 80416 entries, 0 to 80415
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	80416 non-null	int64
1	key	80416 non-null	object
2	fare_amount	80416 non-null	float64
3	pickup_datetime	80416 non-null	object
4	pickup_longitude	80416 non-null	float64
5	pickup_latitude	80416 non-null	float64
6	dropoff_longitude	80416 non-null	float64
7	dropoff_latitude	80415 non-null	float64
8	passenger_count	80415 non-null	float64

```
dtypes: float64(6), int64(1), object(2)
```

```
memory usage: 5.5+ MB
```

```
df.columns #TO get number of columns in the dataset
```

```
Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
       'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
       'dropoff_latitude', 'passenger_count'],
      dtype='object')
```

```
df =df.drop(['Unnamed: 0', 'key'], axis= 1) #To drop unnamed column
```

```
df.head()
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

				pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
				40.738354	-73.999512	40.723217	1.0
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.750325	1.0
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.772647	1.0
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.803349	3.0
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.761247	5.0

df.shape

#To get the total (Rows,Columns)

(80416, 7)

df.dtypes

#To get the type of each column

```
fare_amount      float64
pickup_datetime  object
pickup_longitude float64
pickup_latitude  float64
dropoff_longitude float64
dropoff_latitude float64
passenger_count  float64
dtype: object
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 80416 entries, 0 to 80415
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount            80416 non-null float64
1   pickup_datetime        80416 non-null object
2   pickup_longitude       80416 non-null float64
3   pickup_latitude        80416 non-null float64
4   dropoff_longitude      80416 non-null float64
```

5 dropoff_latitude 80415 non-null float64

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

memory usage: 4.3+ MB

df.describe() #To get statistics of each columns

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
count	80416.000000	80416.000000	80416.000000	80416.000000	80415.000000	80415.000000
mean	11.381542	-72.533096	39.945845	-72.567713	39.934459	1.674874
std	9.924870	11.857315	8.557173	15.738776	6.803074	1.295577
min	-5.000000	-1340.648410	-74.015515	-3356.666300	-74.009465	0.000000
25%	6.000000	-73.992020	40.734812	-73.991417	40.733664	1.000000
50%	8.500000	-73.981775	40.752595	-73.980082	40.752982	1.000000
75%	12.500000	-73.967171	40.767118	-73.963773	40.768112	2.000000
max	350.000000	40.808425	1644.421482	40.828672	872.697628	6.000000



2. Filling Missing values

df.isnull().sum()

```
fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude 1
passenger_count  1
dtype: int64
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```
df['passenger_count'].fillna(value=df['passenger_count']
                             .median(),inplace = True)
df['dropoff_latitude'].fillna(value=df['dropoff_latitude']
                              .mean(),inplace = True)
df['dropoff_longitude'].fillna(value=df['dropoff_longitude']
                               .median(),inplace = True)
```

```
df.isnull().sum()
```

```
fare_amount      0
pickup_datetime  0
pickup_longitude 0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude 0
passenger_count  0
dtype: int64
```

```
df.dtypes
```

```
fare_amount      float64
pickup_datetime  object
pickup_longitude  float64
pickup_latitude  float64
dropoff_longitude float64
dropoff_latitude float64
passenger_count  float64
dtype: object
```

```
df['pickup_datetime']
```

```
0      2015-05-07 19:52:06 UTC
1      2009-07-17 20:04:56 UTC
2      2009-08-24 21:45:00 UTC
3      2009-06-26 08:22:21 UTC
```

```
4      2014-08-28 17:47:00 UTC
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```
80412      2009-04-02 14:35:00 UTC
```

```
80413      2013-05-29 15:25:23 UTC
```

```
80414      2011-11-17 14:49:35 UTC
```

```
80415      2011-11-01 02:33:26 UTC
```

```
Name: pickup_datetime, Length: 80416, dtype: object
```

Date is in incorrct format

```
df.pickup_datetime=pd.to_datetime(df.pickup_datetime,errors='coerce')
```

```
df.dtypes
```

```
fare_amount      float64
```

```
pickup_datetime  datetime64[ns, UTC]
```

```
pickup_longitude float64
```

```
pickup_latitude  float64
```

```
dropoff_longitude float64
```

```
dropoff_latitude  float64
```

```
passenger_count   float64
```

```
dtype: object
```

To seperate each time of date and time

```
df=df.assign(hour = df.pickup_datetime.dt.hour,
day= df.pickup_datetime.dt.day, month = df.pickup_datetime.dt.month,
year = df.pickup_datetime.dt.year,
dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
df.head()
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu

			pickup_longitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month
0	7.5	2015-05-17 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.723217	1.0	19	7
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.750325	1.0	20	17
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.772647	1.0	21	24
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.803349	3.0	8	26
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.761247	5.0	17	28



```
# drop the column 'pickup_datetime' using drop() # 'axis = 1' drops the specified column
df = df.drop('pickup_datetime',axis=1)
```

```
df.head()
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayof
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1.0	19	7	5	2015	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1.0	20	17	7	2009	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1.0	21	24	8	2009	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	3.0	8	26	6	2009	
4	16.0	-73.925023	40.744085	-73.973082	40.761247	5.0	17	28	8	2014	



```
df.dtypes
```

B

```
fare_amount      float64
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```
dropoff_longitude float64
dropoff_latitude   float64
passenger_count     float64
hour                int64
day                 int64
month               int64
year                int64
dayofweek           int64
dtype: object
```

Checking outliers and filling them


```
df.plot(kind = "box",
        subplots = True,
        layout = (4,3),
        figsize=(12,8)) #Boxplot to check
```


To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```

dropoff_longitude    AxesSubplot(0.125,0.518913;0.227941x0.16413)
dropoff_latitude     AxesSubplot(0.398529,0.518913;0.227941x0.16413)
passenger_count      AxesSubplot(0.672059,0.518913;0.227941x0.16413)
hour                 AxesSubplot(0.125,0.321957;0.227941x0.16413)
day                  AxesSubplot(0.398529,0.321957;0.227941x0.16413)
month                AxesSubplot(0.672059,0.321957;0.227941x0.16413)
year                 AxesSubplot(0.125,0.125;0.227941x0.16413)
dayofweek            AxesSubplot(0.398529,0.125;0.227941x0.16413)
dtype: object

```



#Using the InterQuartile Range to fill the values

```

def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df[col] = np.clip(df1[col] ,
                      lower_whisker ,
                      upper_whisker)

    return df1
def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1

df = treat_outliers_all(df , df.iloc[:, 0::])

```



```
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20))
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu

73x0.0920732)
73x0.0920732)
73x0.0920732)

```
dropoff_longitude AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object
```



```
!pip install haversine
import haversine as hs #Calculate the distance using Haversine to calculate the dista
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [
        df['pickup_longitude'][pos],
        df['pickup_latitude'][pos],
        df['dropoff_longitude'][pos],
        df['dropoff_latitude'][pos]
    ]
    loc1=(lati1,long1)
    loc2=(lati2,long2)
    c = hs.haversine(loc1,loc2)
    travel_dist.append(c)
print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu  python.pkg.dev/colab-wheels/public/simple/

3)

Installing collected packages: haversine

Successfully installed haversine-2.7.0

[1.6833250775073447, 2.4575932783467835, 5.036384146783453, 1.661685753650294, 4.107873890221249, 0.0, 9.521855346882292, 0.8032336690

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayof
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1.0	19	7	5	2015	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1.0	20	17	7	2009	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1.0	21	24	8	2009	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	3.0	8	26	6	2009	
4	16.0	-73.929896	40.744085	-73.973082	40.761247	3.5	17	28	8	2014	



```
#Uber doesn't travel over 130 kms so minimize the distance
df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observastions in the dataset:", df.shape)
```

Remaining observastions in the dataset: (80416, 12)

```
#Finding inccorect latitude (Less than or greater than 90) and
#longitude (less than or greater than 180)
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) |
                               (df.pickup_latitude < -90) |
                               (df.dropoff_latitude > 180)|
                               (df.pickup_longitude >-180)|
                               (df.dropoff_longitude > 90)|
                               (df.dropoff_longitude<-90)]
```

```
df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
```

B

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year	dayof
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1.0	19	7	5	2015	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1.0	20	17	7	2009	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1.0	21	24	8	2009	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	3.0	8	26	6	2009	
4	16.0	-73.929896	40.744085	-73.973082	40.761247	3.5	17	28	8	2014	




```
df.isnull().sum()
```

```
fare_amount      0
pickup_longitude  0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude  0
passenger_count   0
hour              0
day               0
month             0
year              0
dayofweek         0
dist_travel_km    0
dtype: int64
```

```
sns.heatmap(df.isnull()) #Free for null values
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕



```
corr =df.corr() #Function to find the correlation  
corr
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu



	latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour
pickup_longitude	0.158149	1.000000	0.259929	0.427068	0.075002
pickup_latitude	-0.113400	0.259929	1.000000	0.049440	0.521595

```
fig,axis = plt.subplots(figsize = (12,8))
#Correlation Heatmap (Light values means highly correlated)
sns.heatmap(df.corr(),annot = True)
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu

Dividing the dataset into feature and target values

```
pickup_longitude - 0.16 1 0.26 0.43 0.075 -0.0097 0.0089 -0.0045 0.0019 0.012 -0.028 0.049
df.columns

Index(['fare_amount', 'pickup_longitude', 'pickup_latitude',
      'dropoff_longitude', 'dropoff_latitude', 'passenger_count', 'hour',
      'day', 'month', 'year', 'dayofweek', 'dist_travel_km'],
      dtype='object')

x=df[['pickup_longitude','pickup_latitude','dropoff_longitude',
      'dropoff_latitude','passenger_count','hour','day','month',
      'year','dayofweek','dist_travel_km']]

y=df[['fare_amount']]
```

Dividing the dataset into training and testing dataset

```
dayofweek - 0.0098 -0.028 -0.042 -0.003 -0.031 0.048 -0.031 0.0037 -0.0005 0.0038 1 0.027
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.20)
```

Linear Regression

```
from sklearn.linear_model import LinearRegression
regression = LinearRegression()

regression.fit(X_train,y_train)

LinearRegression()

regression.intercept_ #To find the linear intercept
```


To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```
regression.coef_ #To find the linear coeeficient
```

```
array([[ 2.57457893e+01, -6.61066060e+00,  2.10215581e+01,  
       -1.93587498e+01,  5.38030754e-02,  8.88442096e-03,  
         4.29939852e-03,  5.92900691e-02,  3.69521146e-01,  
       -3.55853647e-02,  1.84826576e+00]])
```

```
prediction = regression.predict(X_test) #To predict the target values
```

```
print(prediction)
```

```
[[17.50994121]  
 [ 6.99860088]  
 [10.84292277]  
 ...  
 [10.72940609]  
 [ 6.54627526]  
 [ 6.36345303]]
```

```
y_test
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

27915 4.9

50169 17.0

41698 7.5

```
from sklearn.metrics import r2_score
```

```
r2_score(y_test,prediction)
```

```
0.682803439471712
```

```
from sklearn.metrics import mean_squared_error
```

11670 7.1

```
MSE=mean_squared_error(y_test,prediction)
```

```
MSE
```

```
9.484117197738067
```

```
RMSE = np.sqrt(MSE)
```

```
RMSE
```

```
3.0796293929202045
```

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

```
rf =RandomForestRegressor(n_estimators=100) #Here n_estimators means number of trees
```

```
rf.fit(X_train,y_train)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please use the ndarray array equivalent y.ravel() to make this warning silent and raise it as a ValueError in an explicit assertion in _fit.  
"""Entry point for launching an IPython kernel.  
RandomForestRegressor()
```

```
y_pred = rf.predict(X_test)
```

```
y_pred
```

```
array([20.0525,  7.352 , 12.4283, ..., 10.2945,  6.366 ,  4.515 ])
```

Metrics evaluatin for Random Forest

```
R2_Random = r2_score(y_test,y_pred)
```

```
R2_Random
```

```
0.7968363495400521
```

```
MSE_Random = mean_squared_error(y_test,y_pred)
```

```
MSE_Random
```

```
6.0745547431955815
```

B

+ Code + Text

To undo cell deletion use Ctrl+M Z or the Undo option in the Edit menu ✕

RMSE_Random

2.464661182230852

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 11:11 AM

